

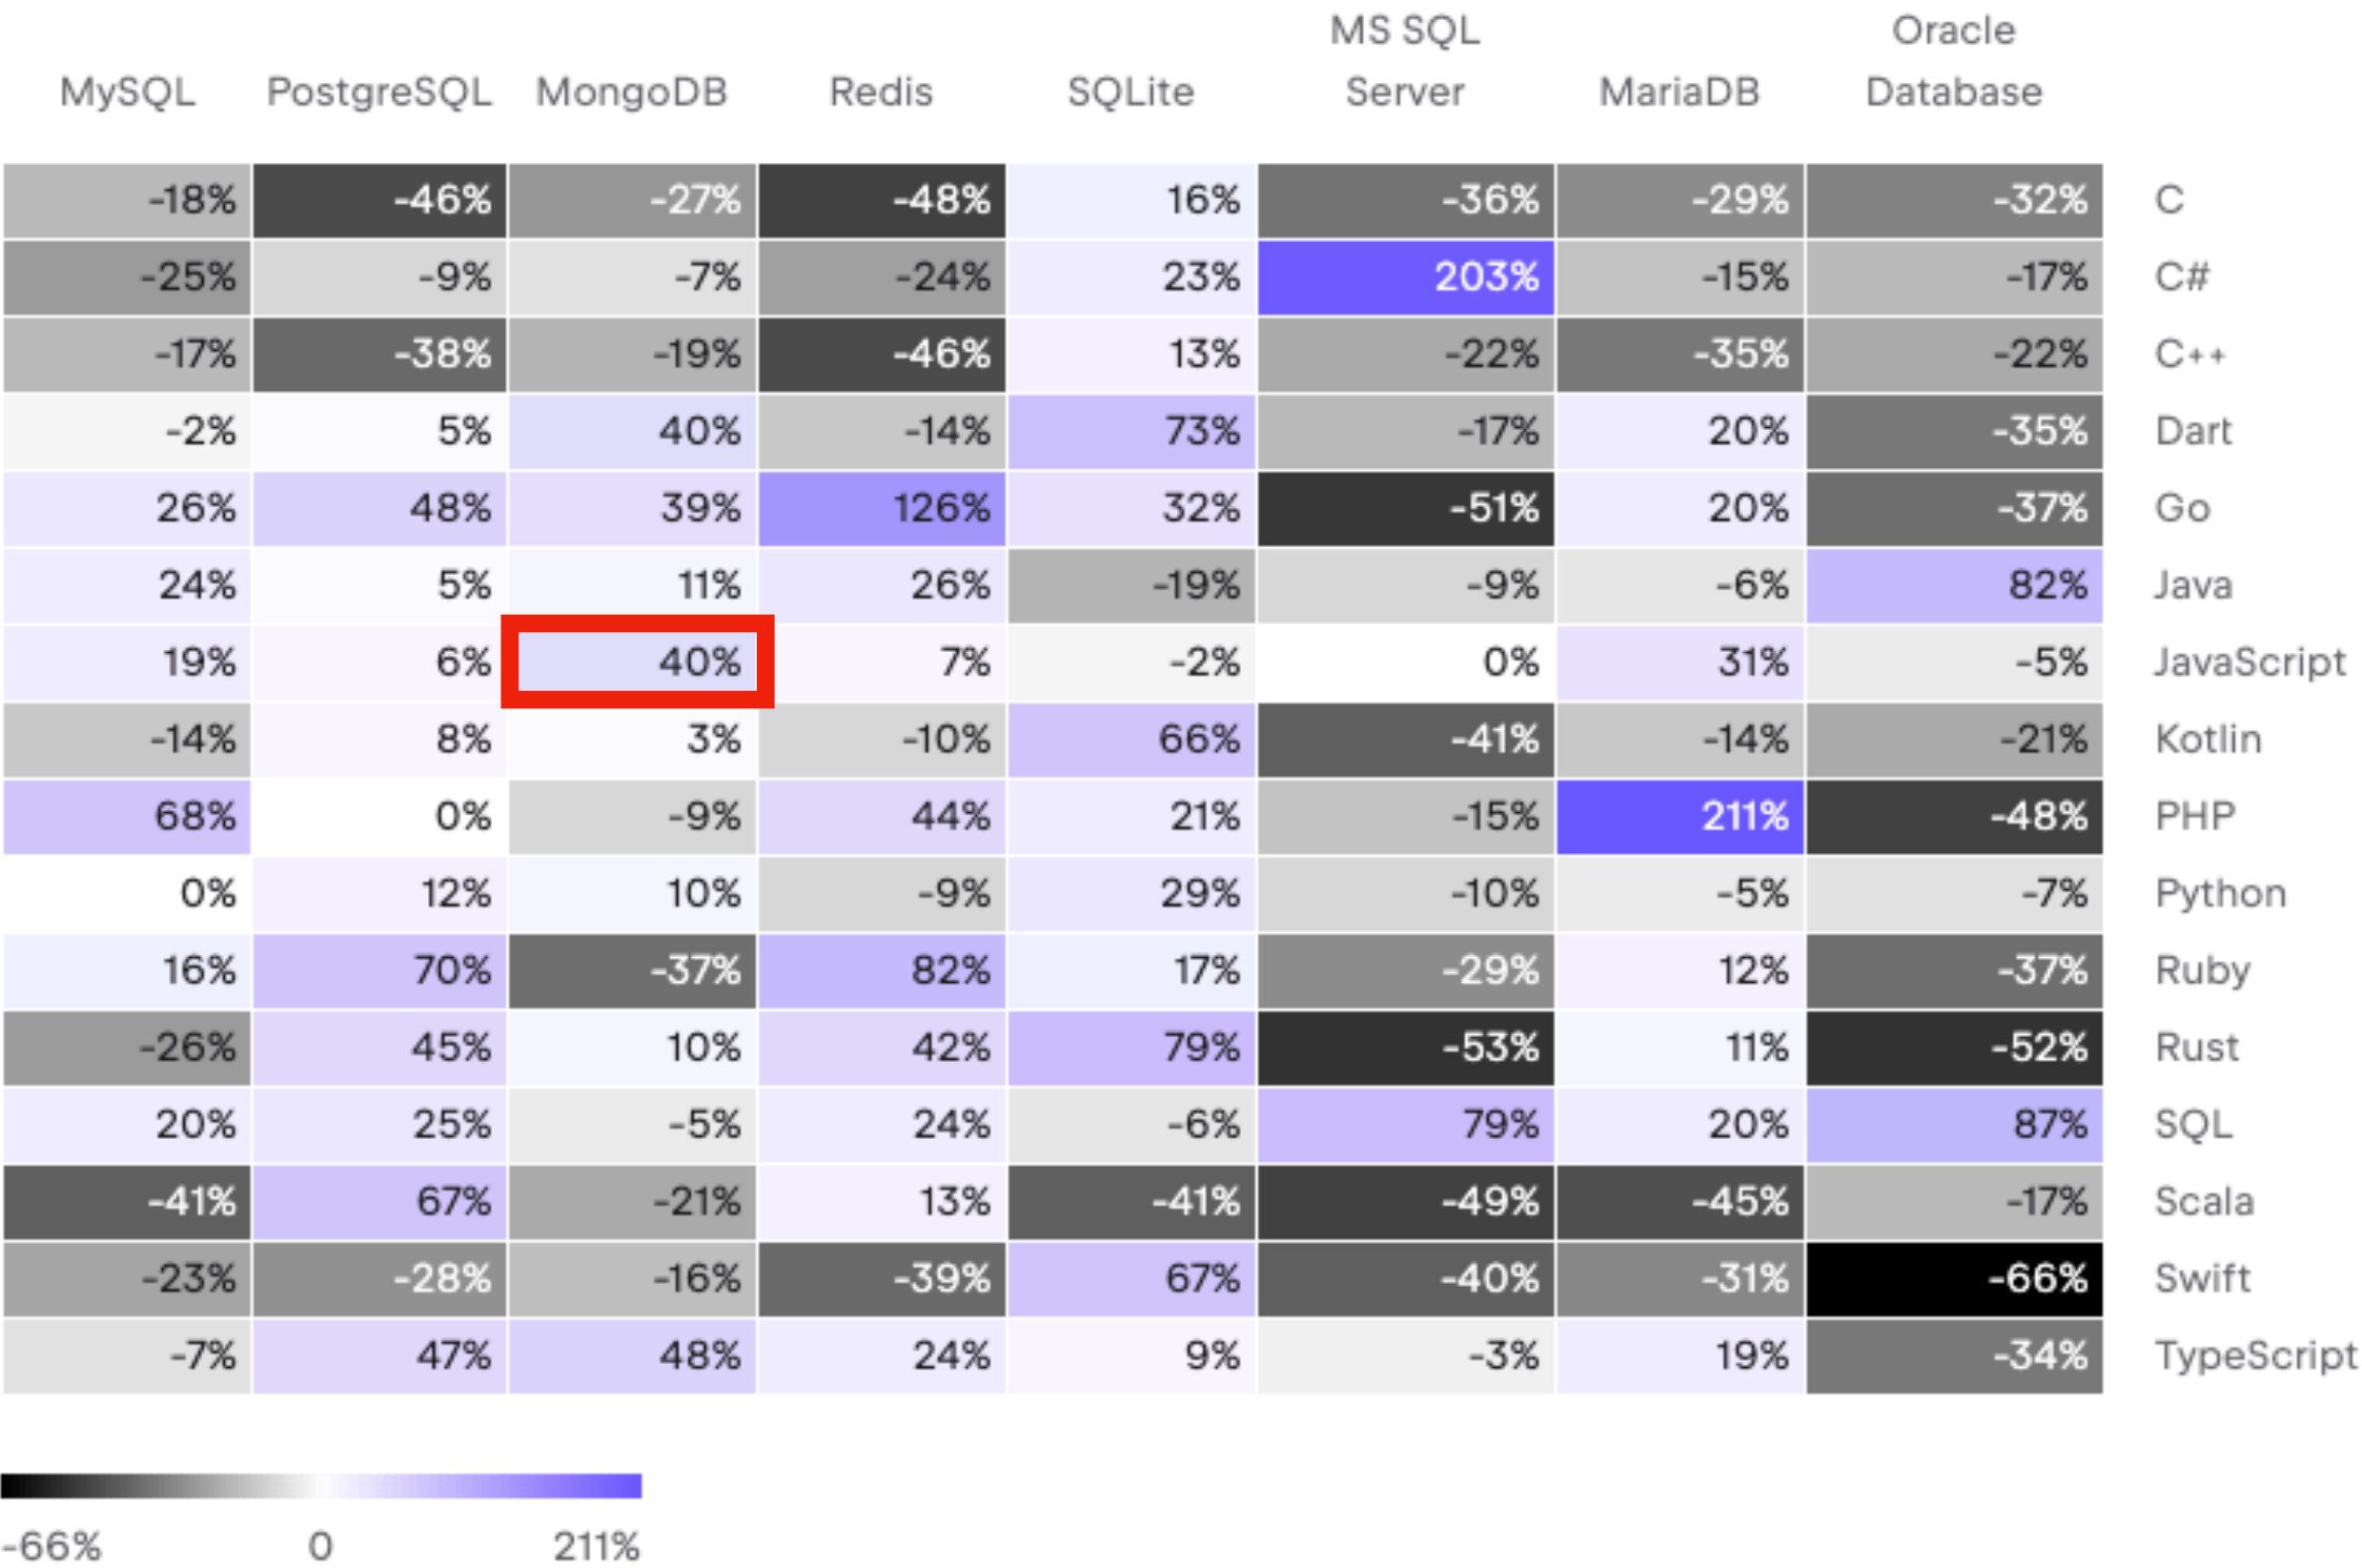

Introducing MongoDB

NoSQL og DBMS



















Which databases have you used in the last 12 months?

2022

2023



423 systems in ranking, October 2024

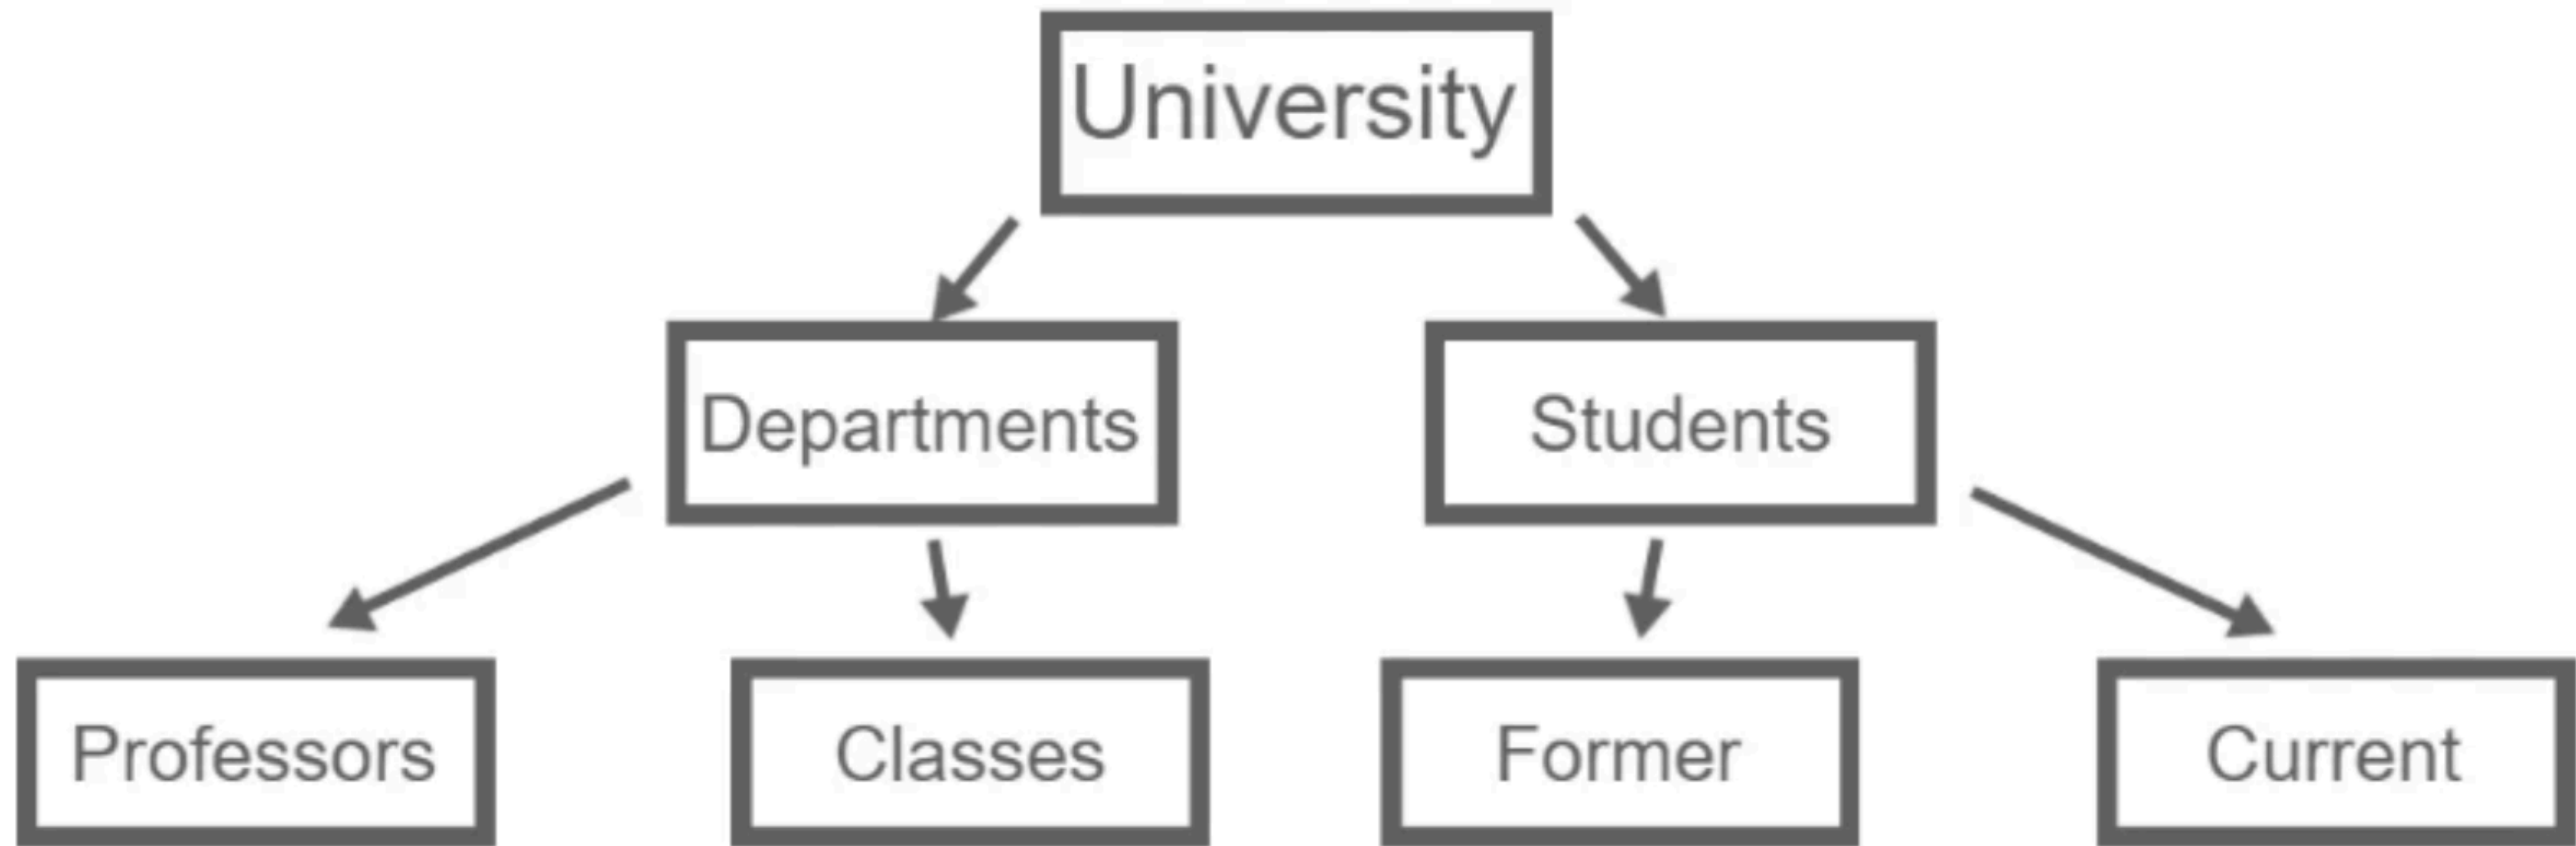
Rank			DBMS	Database Model	Score		
Oct 2024	Sep 2024	Oct 2023			Oct 2024	Sep 2024	Oct 2023
1.	1.	1.	Oracle 	Relational, Multi-model 	1309.45	+22.85	+48.03
2.	2.	2.	MySQL 	Relational, Multi-model 	1022.76	-6.73	-110.56
3.	3.	3.	Microsoft SQL Server	Relational, Multi-model 	802.09	-5.67	-94.79
4.	4.	4.	PostgreSQL 	Relational, Multi-model 	652.16	+7.80	+13.34
5.	5.	5.	MongoDB 	Document, Multi-model 	405.21	-5.02	-26.21
6.	6.	6.	Redis 	Key-value, Multi-model 	149.63	+0.20	-13.33
7.	7.	 11.	Snowflake 	Relational	140.60	+6.88	+17.36
8.	8.	 7.	Elasticsearch	Multi-model 	131.85	+3.06	-5.30
9.	9.	 8.	IBM Db2	Relational, Multi-model 	122.77	-0.28	-12.10
10.	10.	 9.	SQLite	Relational	101.91	-1.43	-23.23

Formål og læringsmål

Fagelementet databasedesign indeholder hvordan man ud fra en datamodel implementerer og vedligeholder en relationel database og sætter data ind i denne på en sikker måde. Derudover indeholder fagelementet alternativer til relationelle databaser og analyseret hvilken databaseteknologi der passer bedst til en given opgave.

Hierarchical Data Model 1960's

(Top to bottom / One to many)



Abstraction: Data is a hierarchy of connected entities from parent to child

```

RD  sales-report
    PAGE LIMITS 60 LINES
    FIRST DETAIL 3
    CONTROLS seller-name.

01  TYPE PAGE HEADING.
    03  COL 1          VALUE "Sales Report".
    03  COL 74         VALUE "Page".
    03  COL 79         PIC Z9 SOURCE PAGE-COUNTER.

01  sales-on-day TYPE DETAIL, LINE + 1.
    03  COL 3          VALUE "Sales on".
    03  COL 12         PIC 99/99/9999 SOURCE sales-date.
    03  COL 21         VALUE "were".
    03  COL 26         PIC $$$9.99 SOURCE sales-amount.

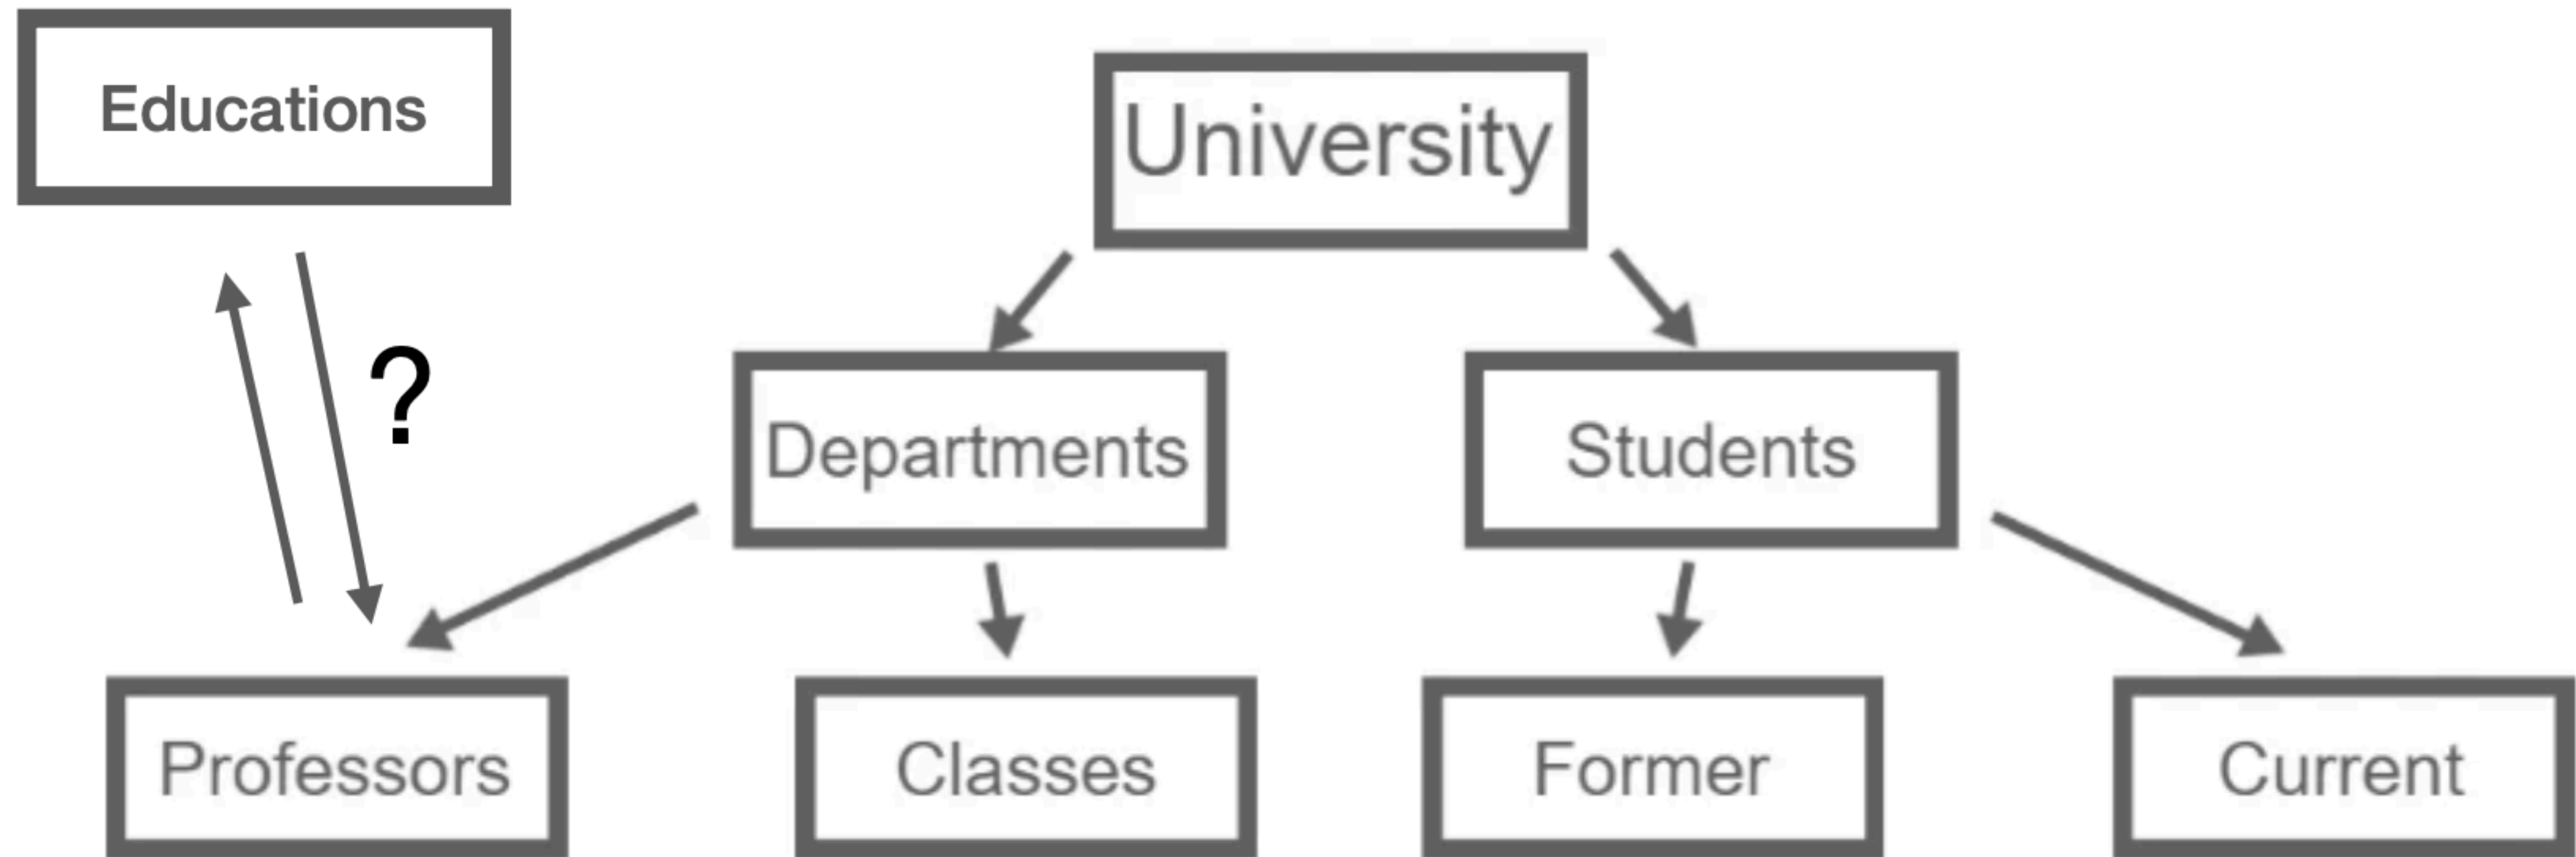
01  invalid-sales TYPE DETAIL, LINE + 1.
    03  COL 3          VALUE "INVALID RECORD:".
    03  COL 19         PIC X(34) SOURCE sales-record.

01  TYPE CONTROL HEADING seller-name, LINE + 2.
    03  COL 1          VALUE "Seller:".
    03  COL 9          PIC X(30) SOURCE seller-name.

```

COBOL language

Hierarchical Data Model 1960's



The relational model 1970

2.4. SUMMARY

In Section 1 a relational model of data is proposed as a basis for protecting users of formatted data systems from the potentially disruptive changes in data representation caused by growth in the data bank and changes in traffic.

In Section 2 operations on relations and two types of redundancy are defined and applied to the problem of maintaining the data in a consistent state. This is bound to become a serious practical problem as more and more different types of data are integrated together into common data banks.

Existing non-inferential, formatted data systems provide users with tree-structured files or slightly more general network models of the data. In Section 1, inadequacies of these models are discussed. A model based on n -ary relations, a normal form for data base relations, and the concept of a universal data sublanguage are introduced. In Section 2, certain operations on relations (other than logical inference) are discussed and applied to the problems of redundancy and consistency in the user's model.

TL:DR

E.F Codd's analysis

- Tree or network models are inadequate
- Data models are often subject to change
- A model based on relations offers more flexibility
- In addition - a common query language would be practical (SQL)

Abstraction: Data is a set of entities with interconnected relationships

What is the point?


Brief history of databases

- **Abstractions** are the essential characteristics of a system
- **Abstractions** lies at the heart of any **implementation** of a DBMS
- Technologies represents several implementations of a data storage **abstraction**

- RDBMS:



Relational Databases are (still) dominating the landscape



Supplementary notes

According to the source, the [ranking](#) is a measure of the following parameters:

- website mentions;
- search frequency;
- technical discussion frequency;
- current job offers;
- professional network profiles;
- social network relevance.

 RDBMS

The Cap Theorem

What do we want from a database?



Consistency

Availability


Partition Tolerance

The Cap Theorem

Partitioning

Horizontal Scaling

(Add more same-size nodes)

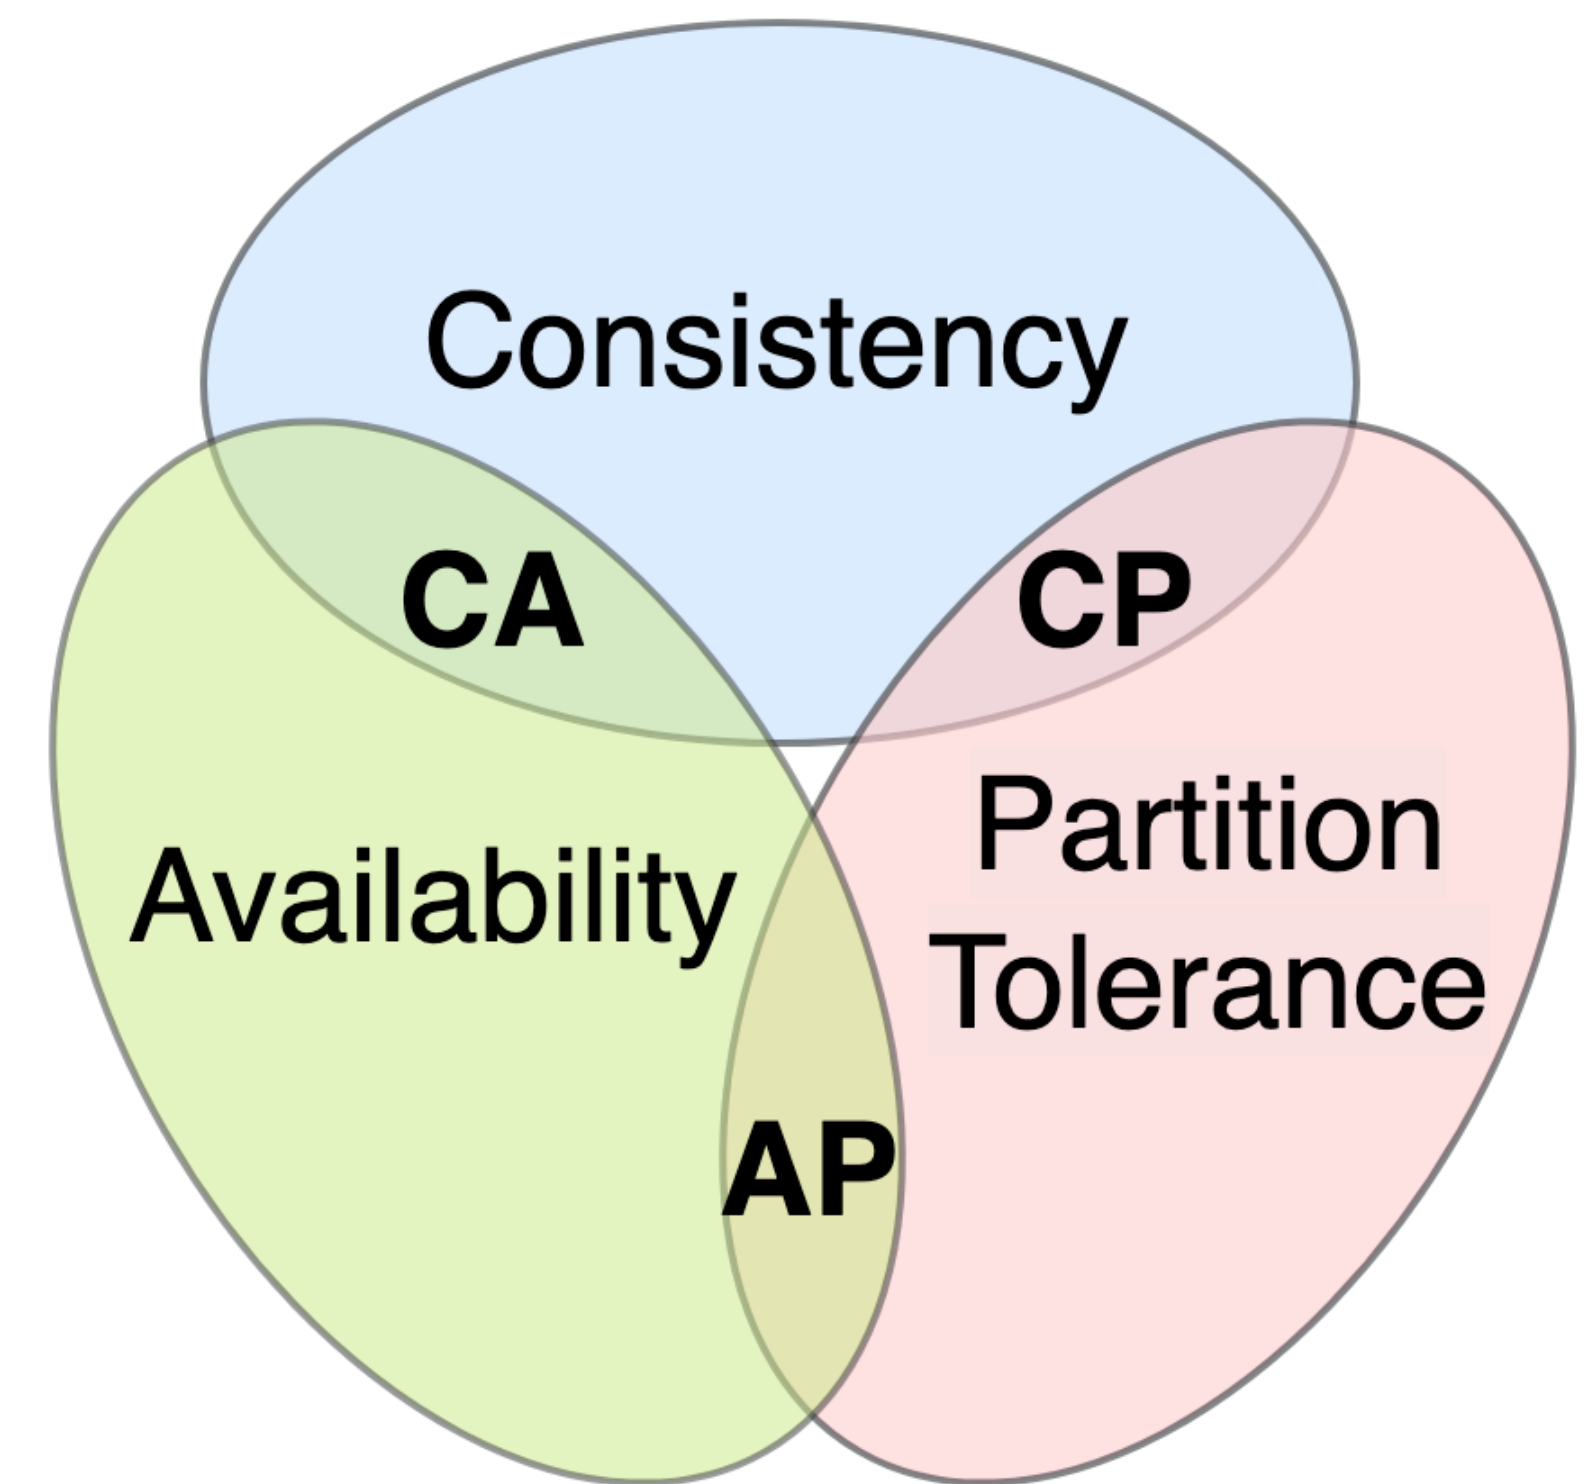


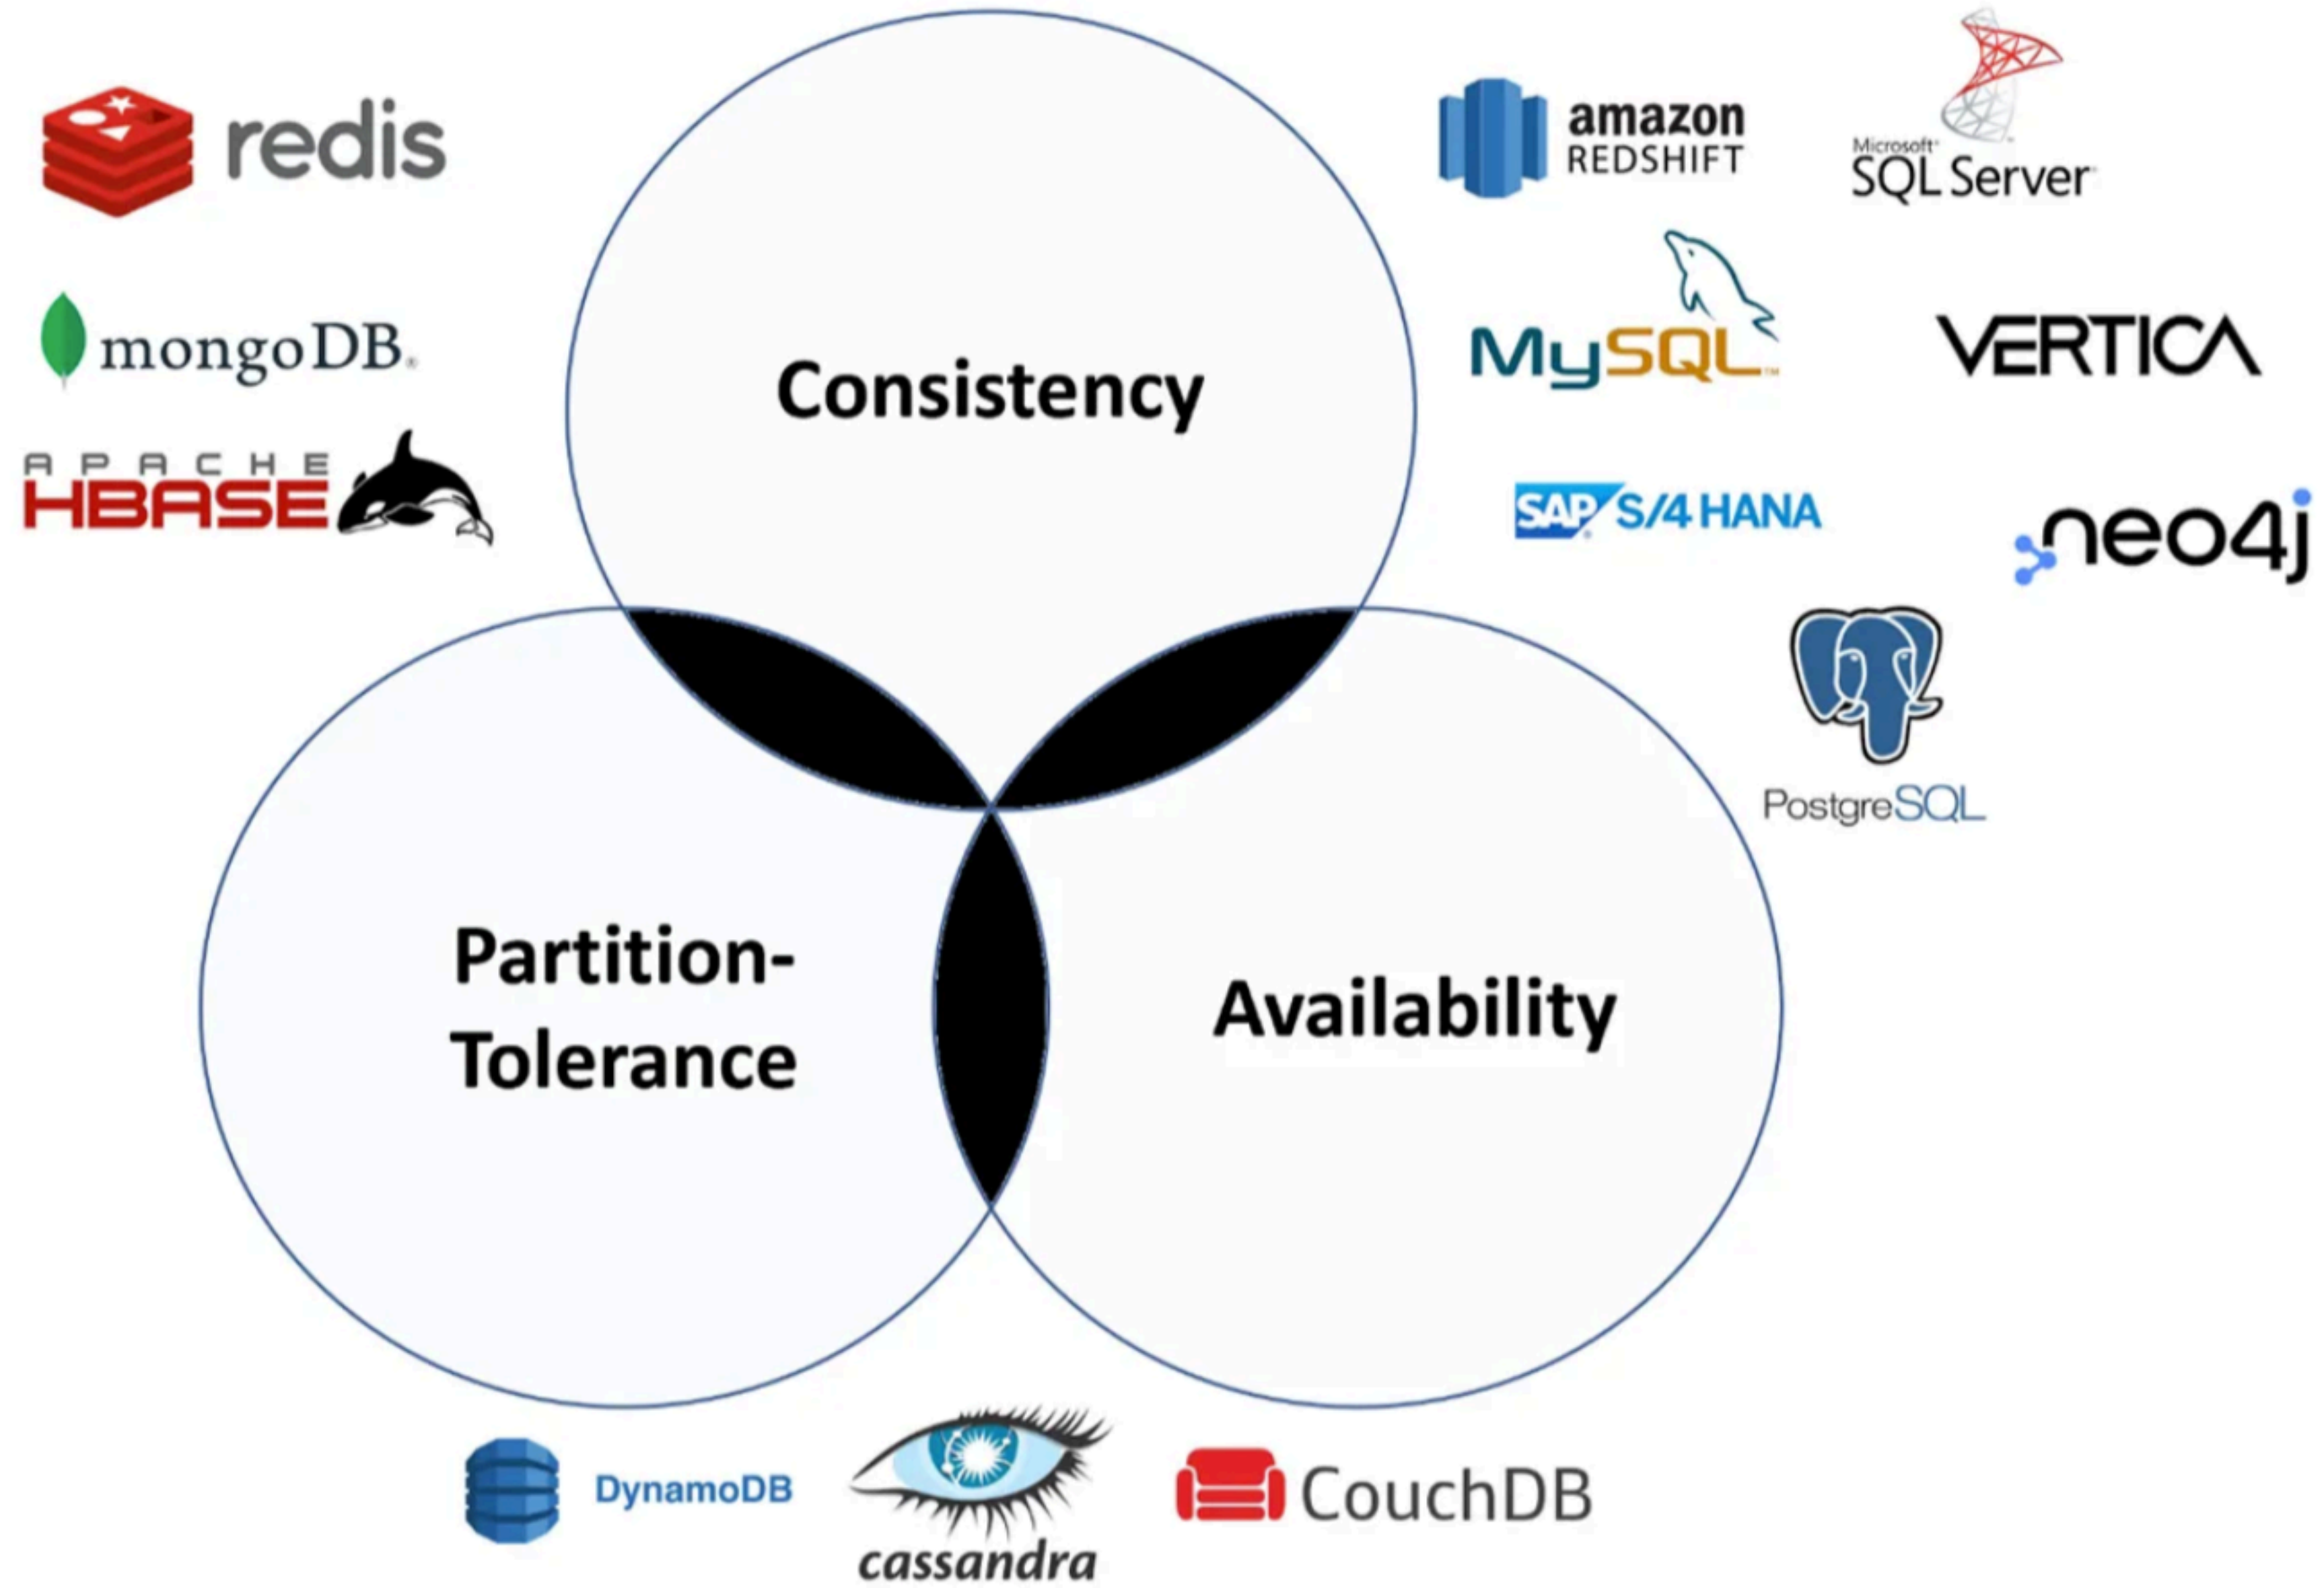
Data is partitioned

The Cap Theorem

Pick 2 - not 3

- **Consistency:** Every read reflects the most recent write. When you update some data, every user or system trying to access data should see the latest version
- **Availability:** Every request gets a response (success or failure), but it does not guarantee that the data is up-to-date.
- **Partition tolerance:** The database continues to operate even if communication between parts of the system is lost (if the network fails)

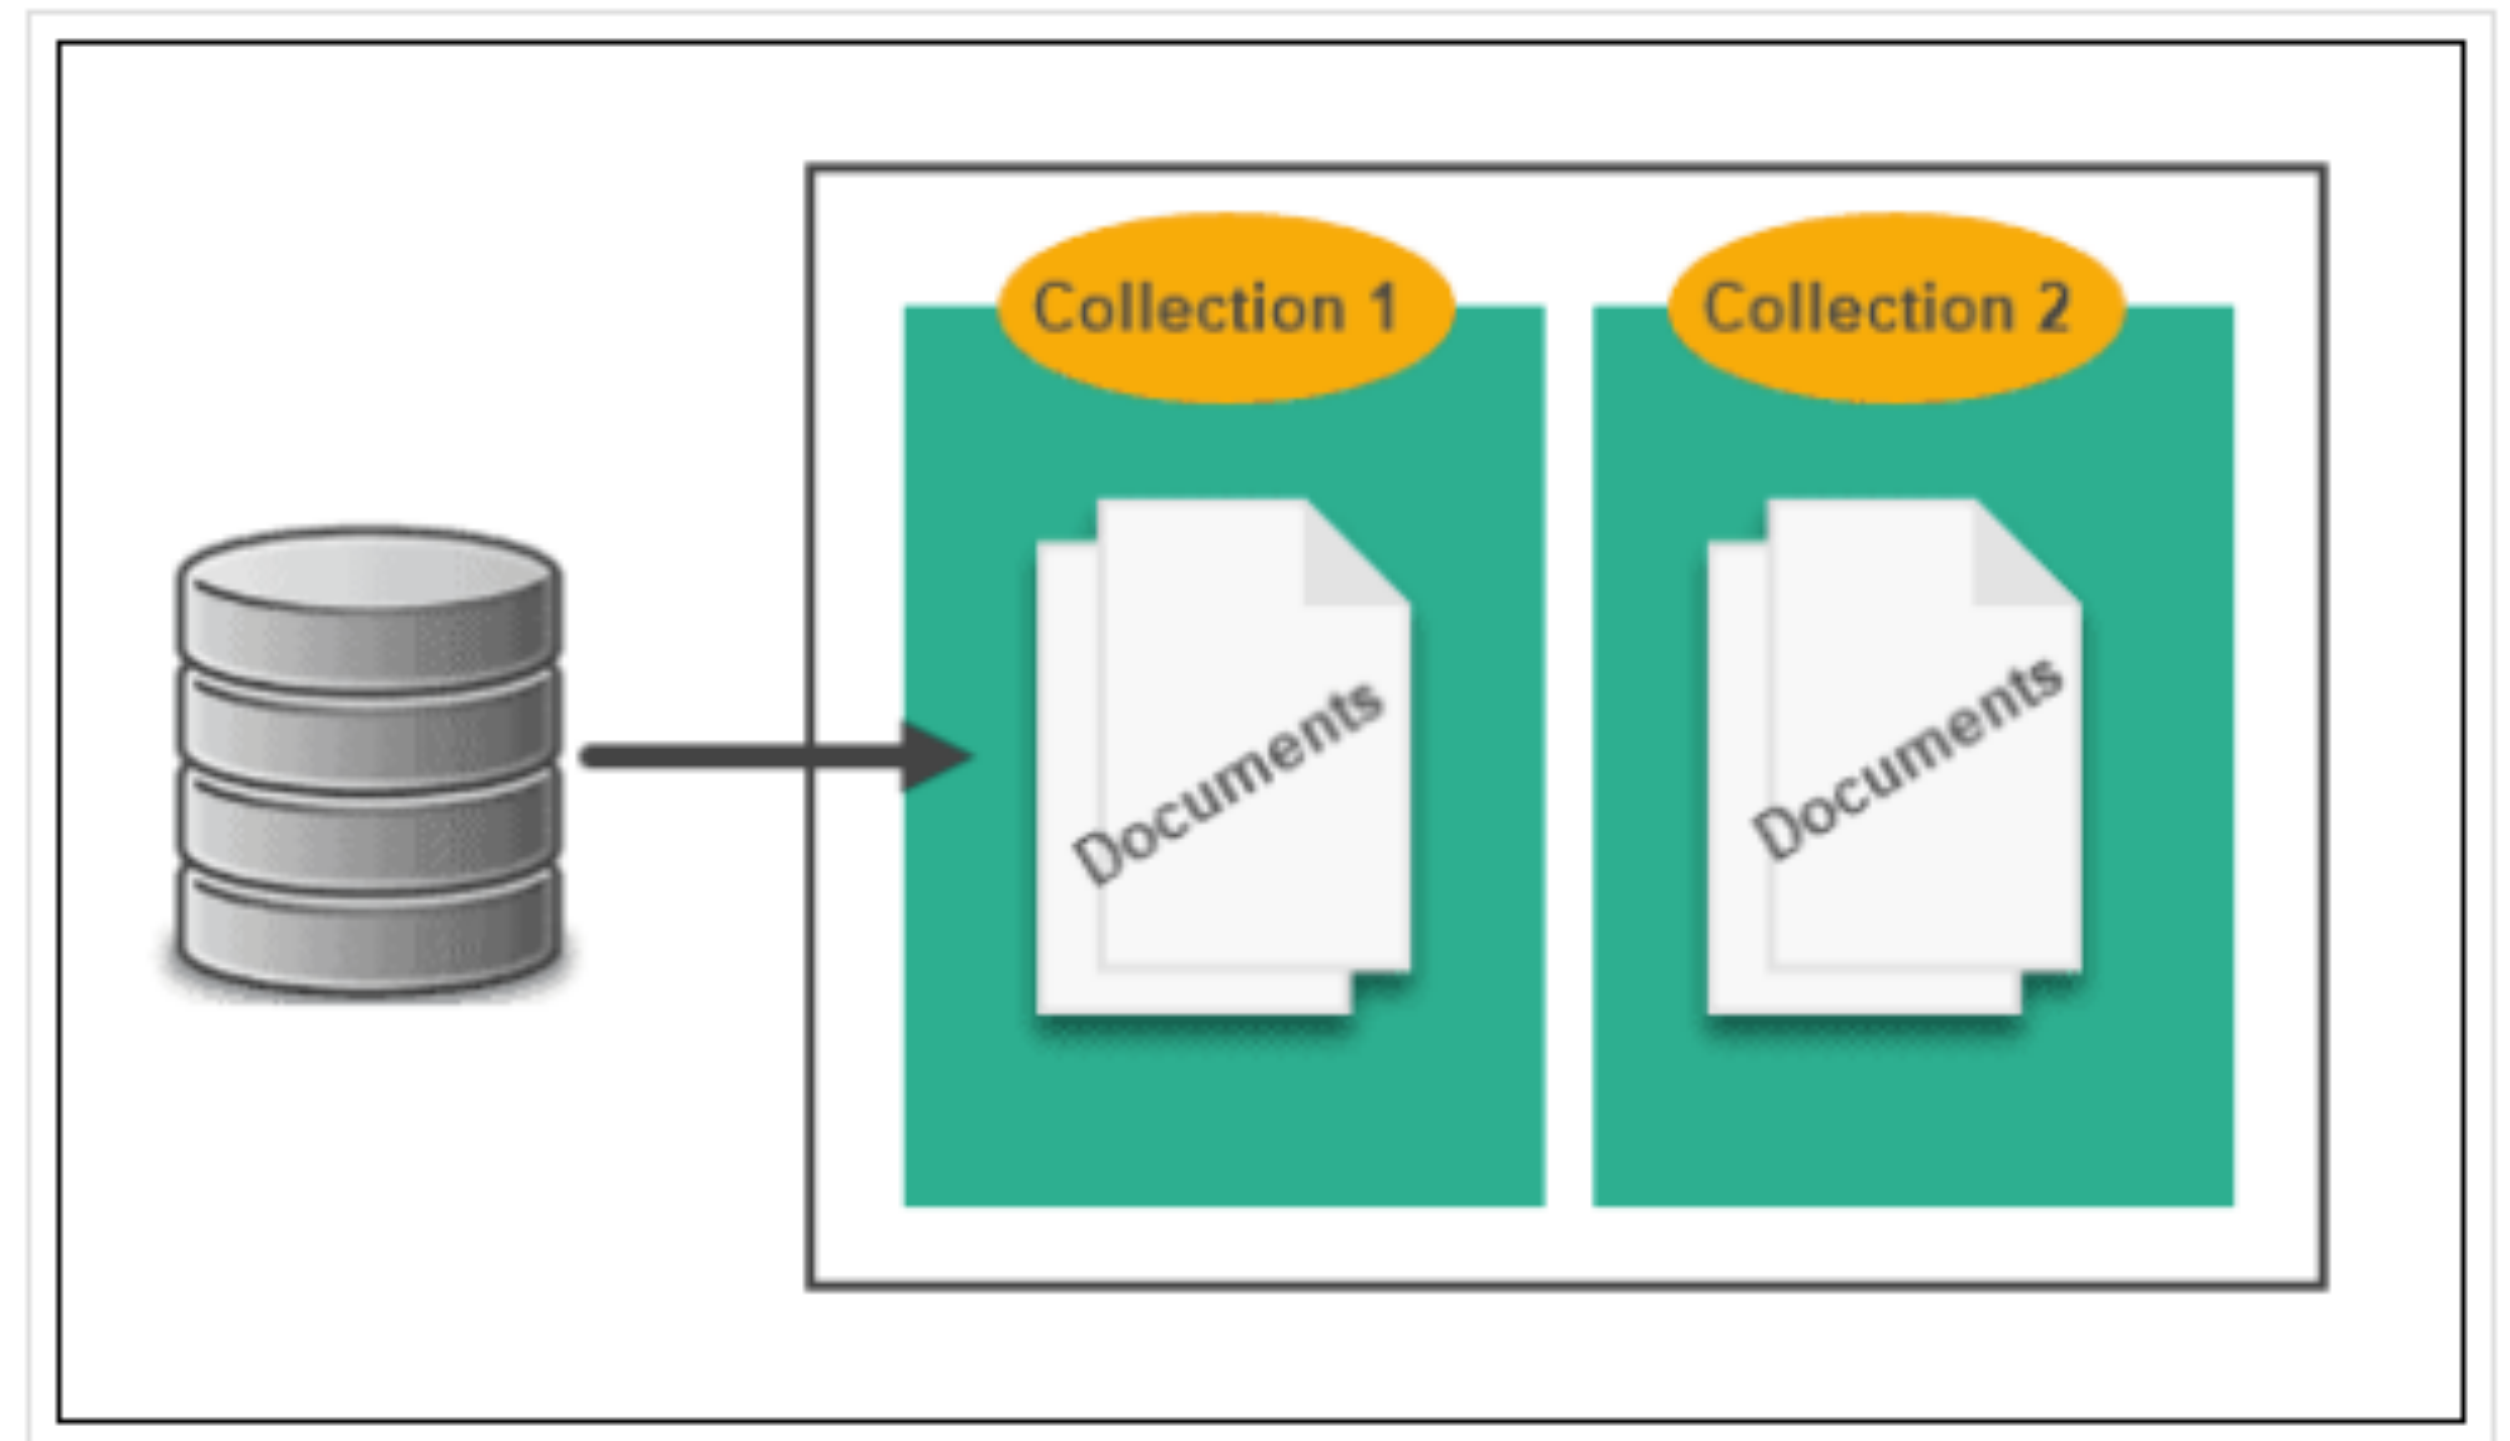




NoSQL

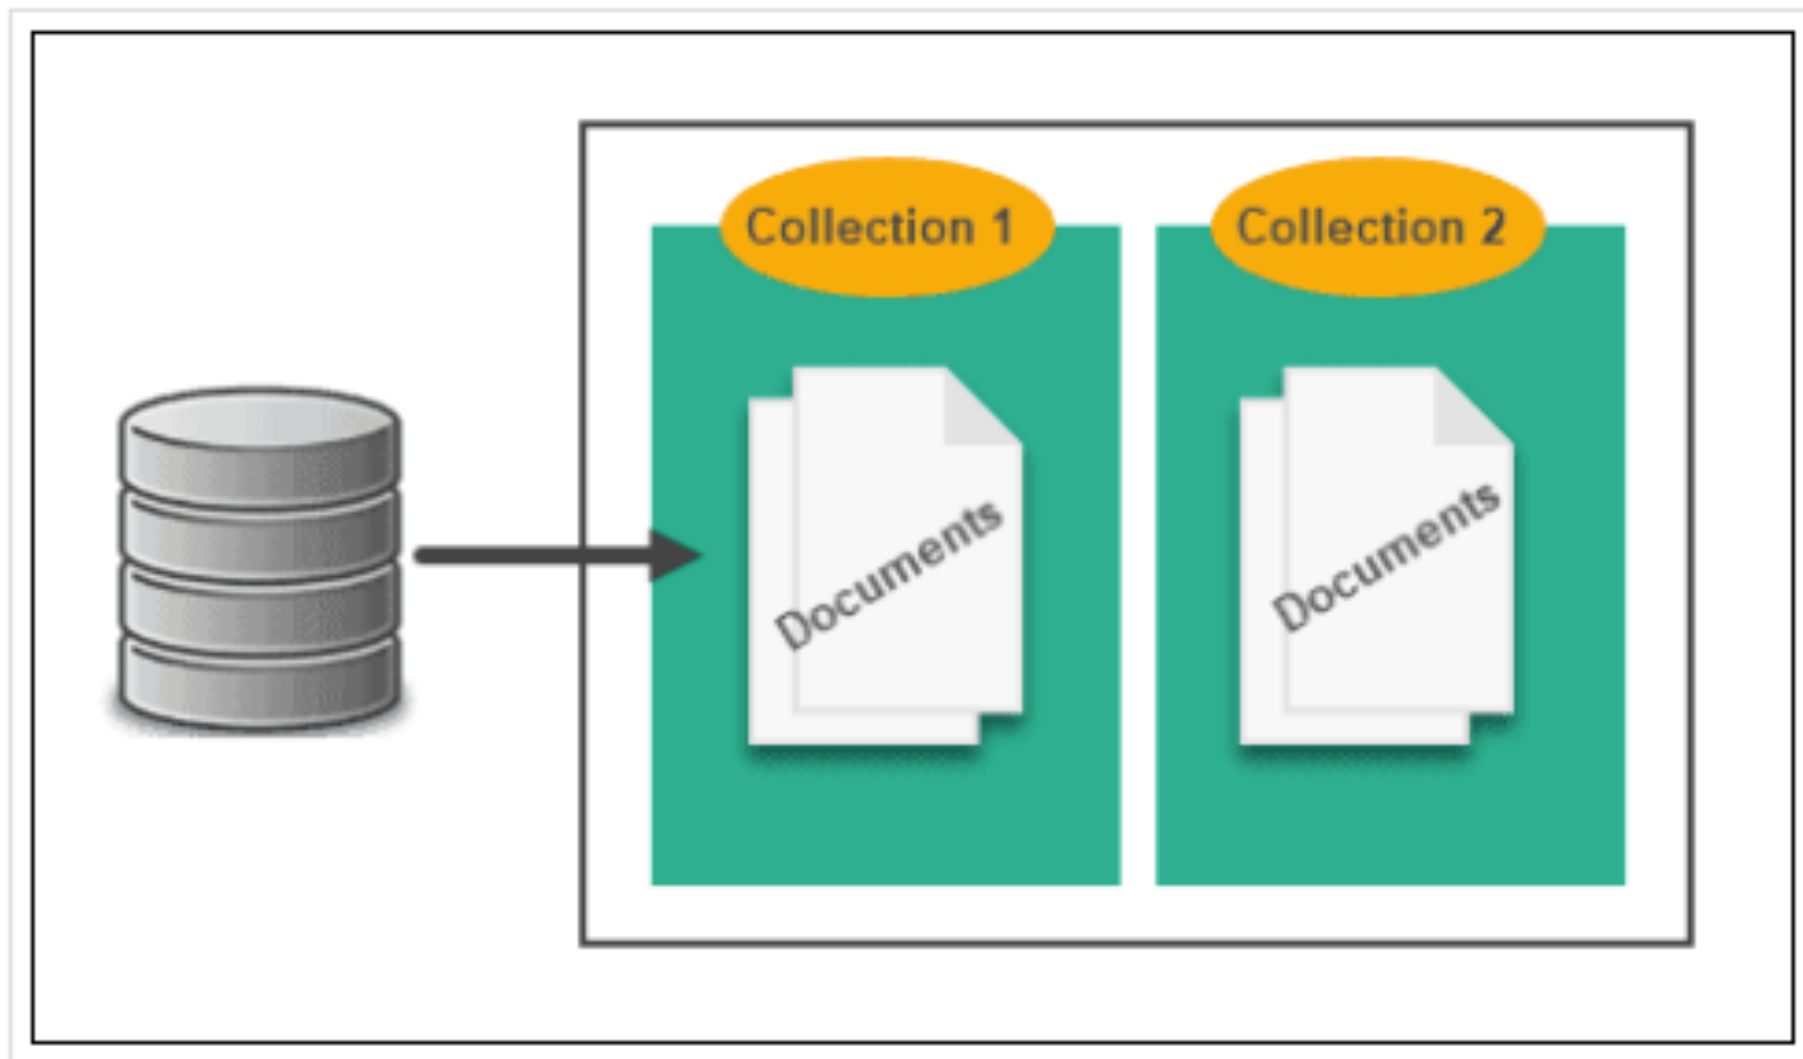
Database types

- Document databases
- Key-Value store databases
- Wide-column stores
- Graph databases



NoSQL

Document database abstraction




```
1  {
2    "_id": "12345",
3    "name": "foo bar",
4    "email": "foo@bar.com",
5    "address": {
6      "street": "123 foo street",
7      "city": "some city",
8      "state": "some state",
9      "zip": "123456"
10   },
11   "hobbies": ["music", "guitar", "reading"]
12 }
```

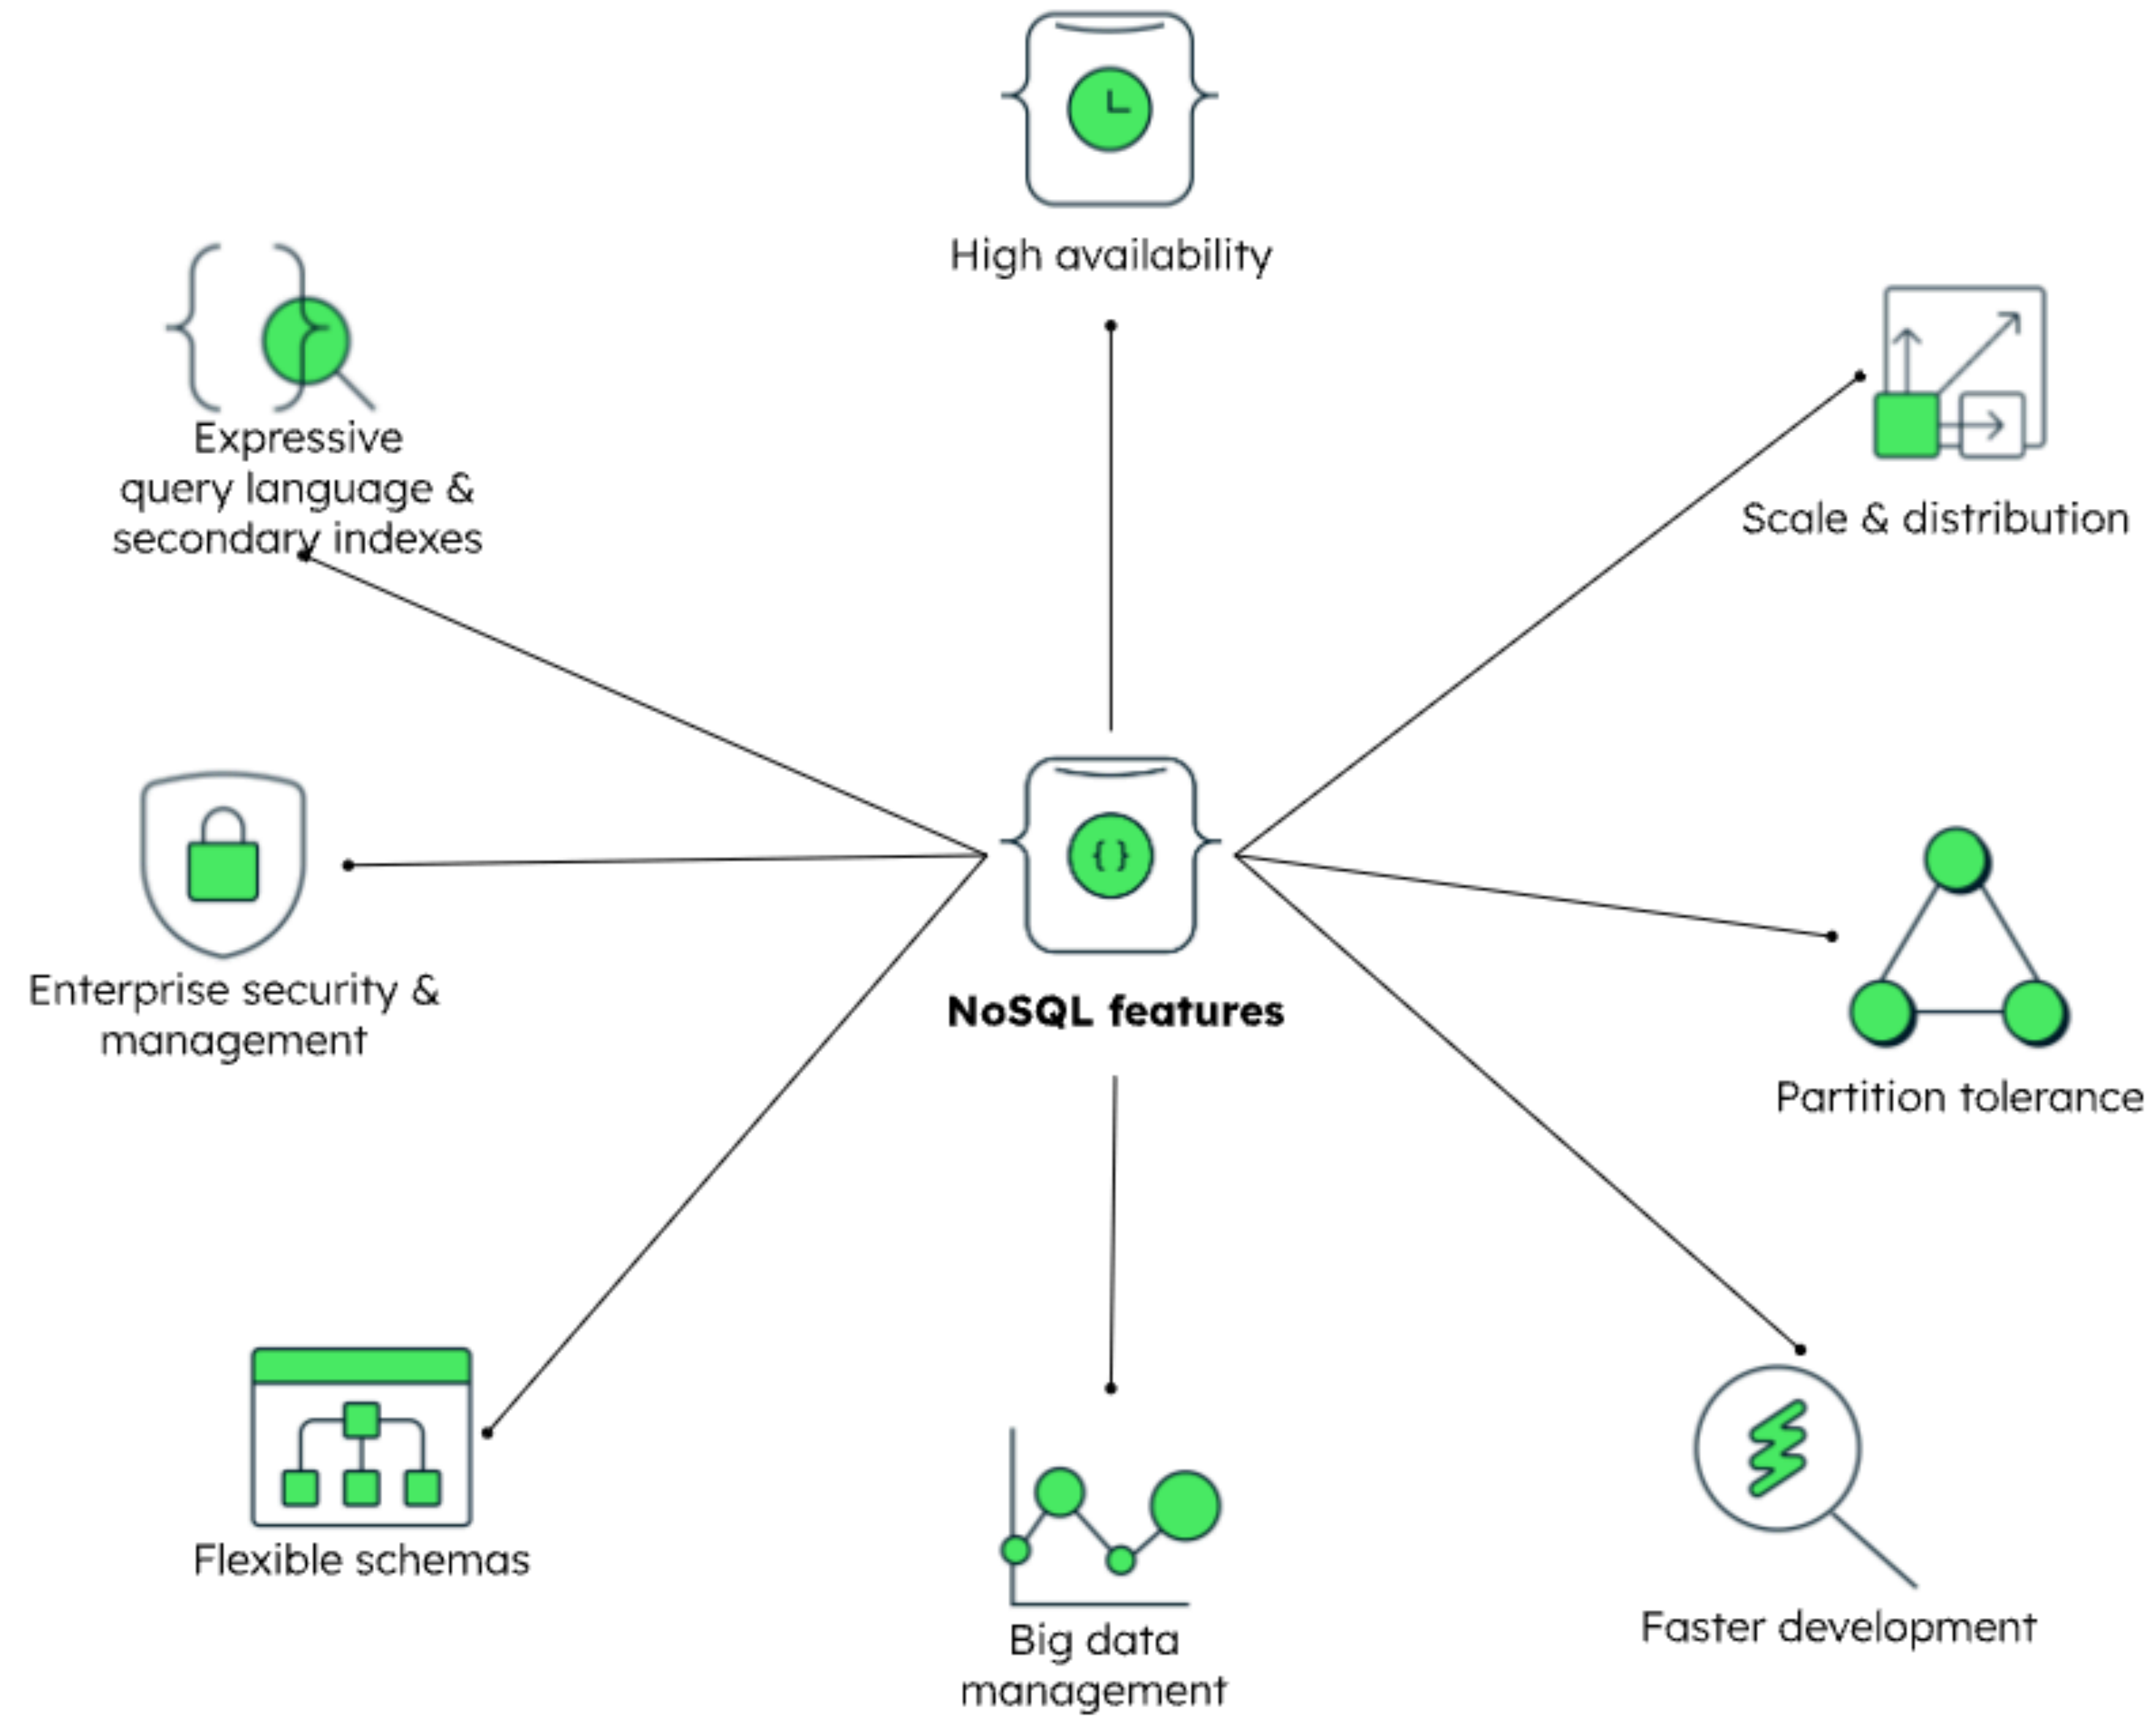


NoSQL


Relevance & emergence

- NoSQL databases emerged in the late 2000s as the cost of storage dramatically decreased. Gone were the days of needing to create a complex, difficult-to-manage data model in order to avoid data duplication. NoSQL databases optimized for developer productivity.





RDBMS vs NoSQL (Document)



Relational
Database

User table

ID	first_name	last_name	cell	city
1	Leslie	Yepp	8125552344	Pawnee

Hobbies table

ID	user_id	hobby
10	1	scrapbooking
11	1	eating waffles
12	1	working



MongoDB

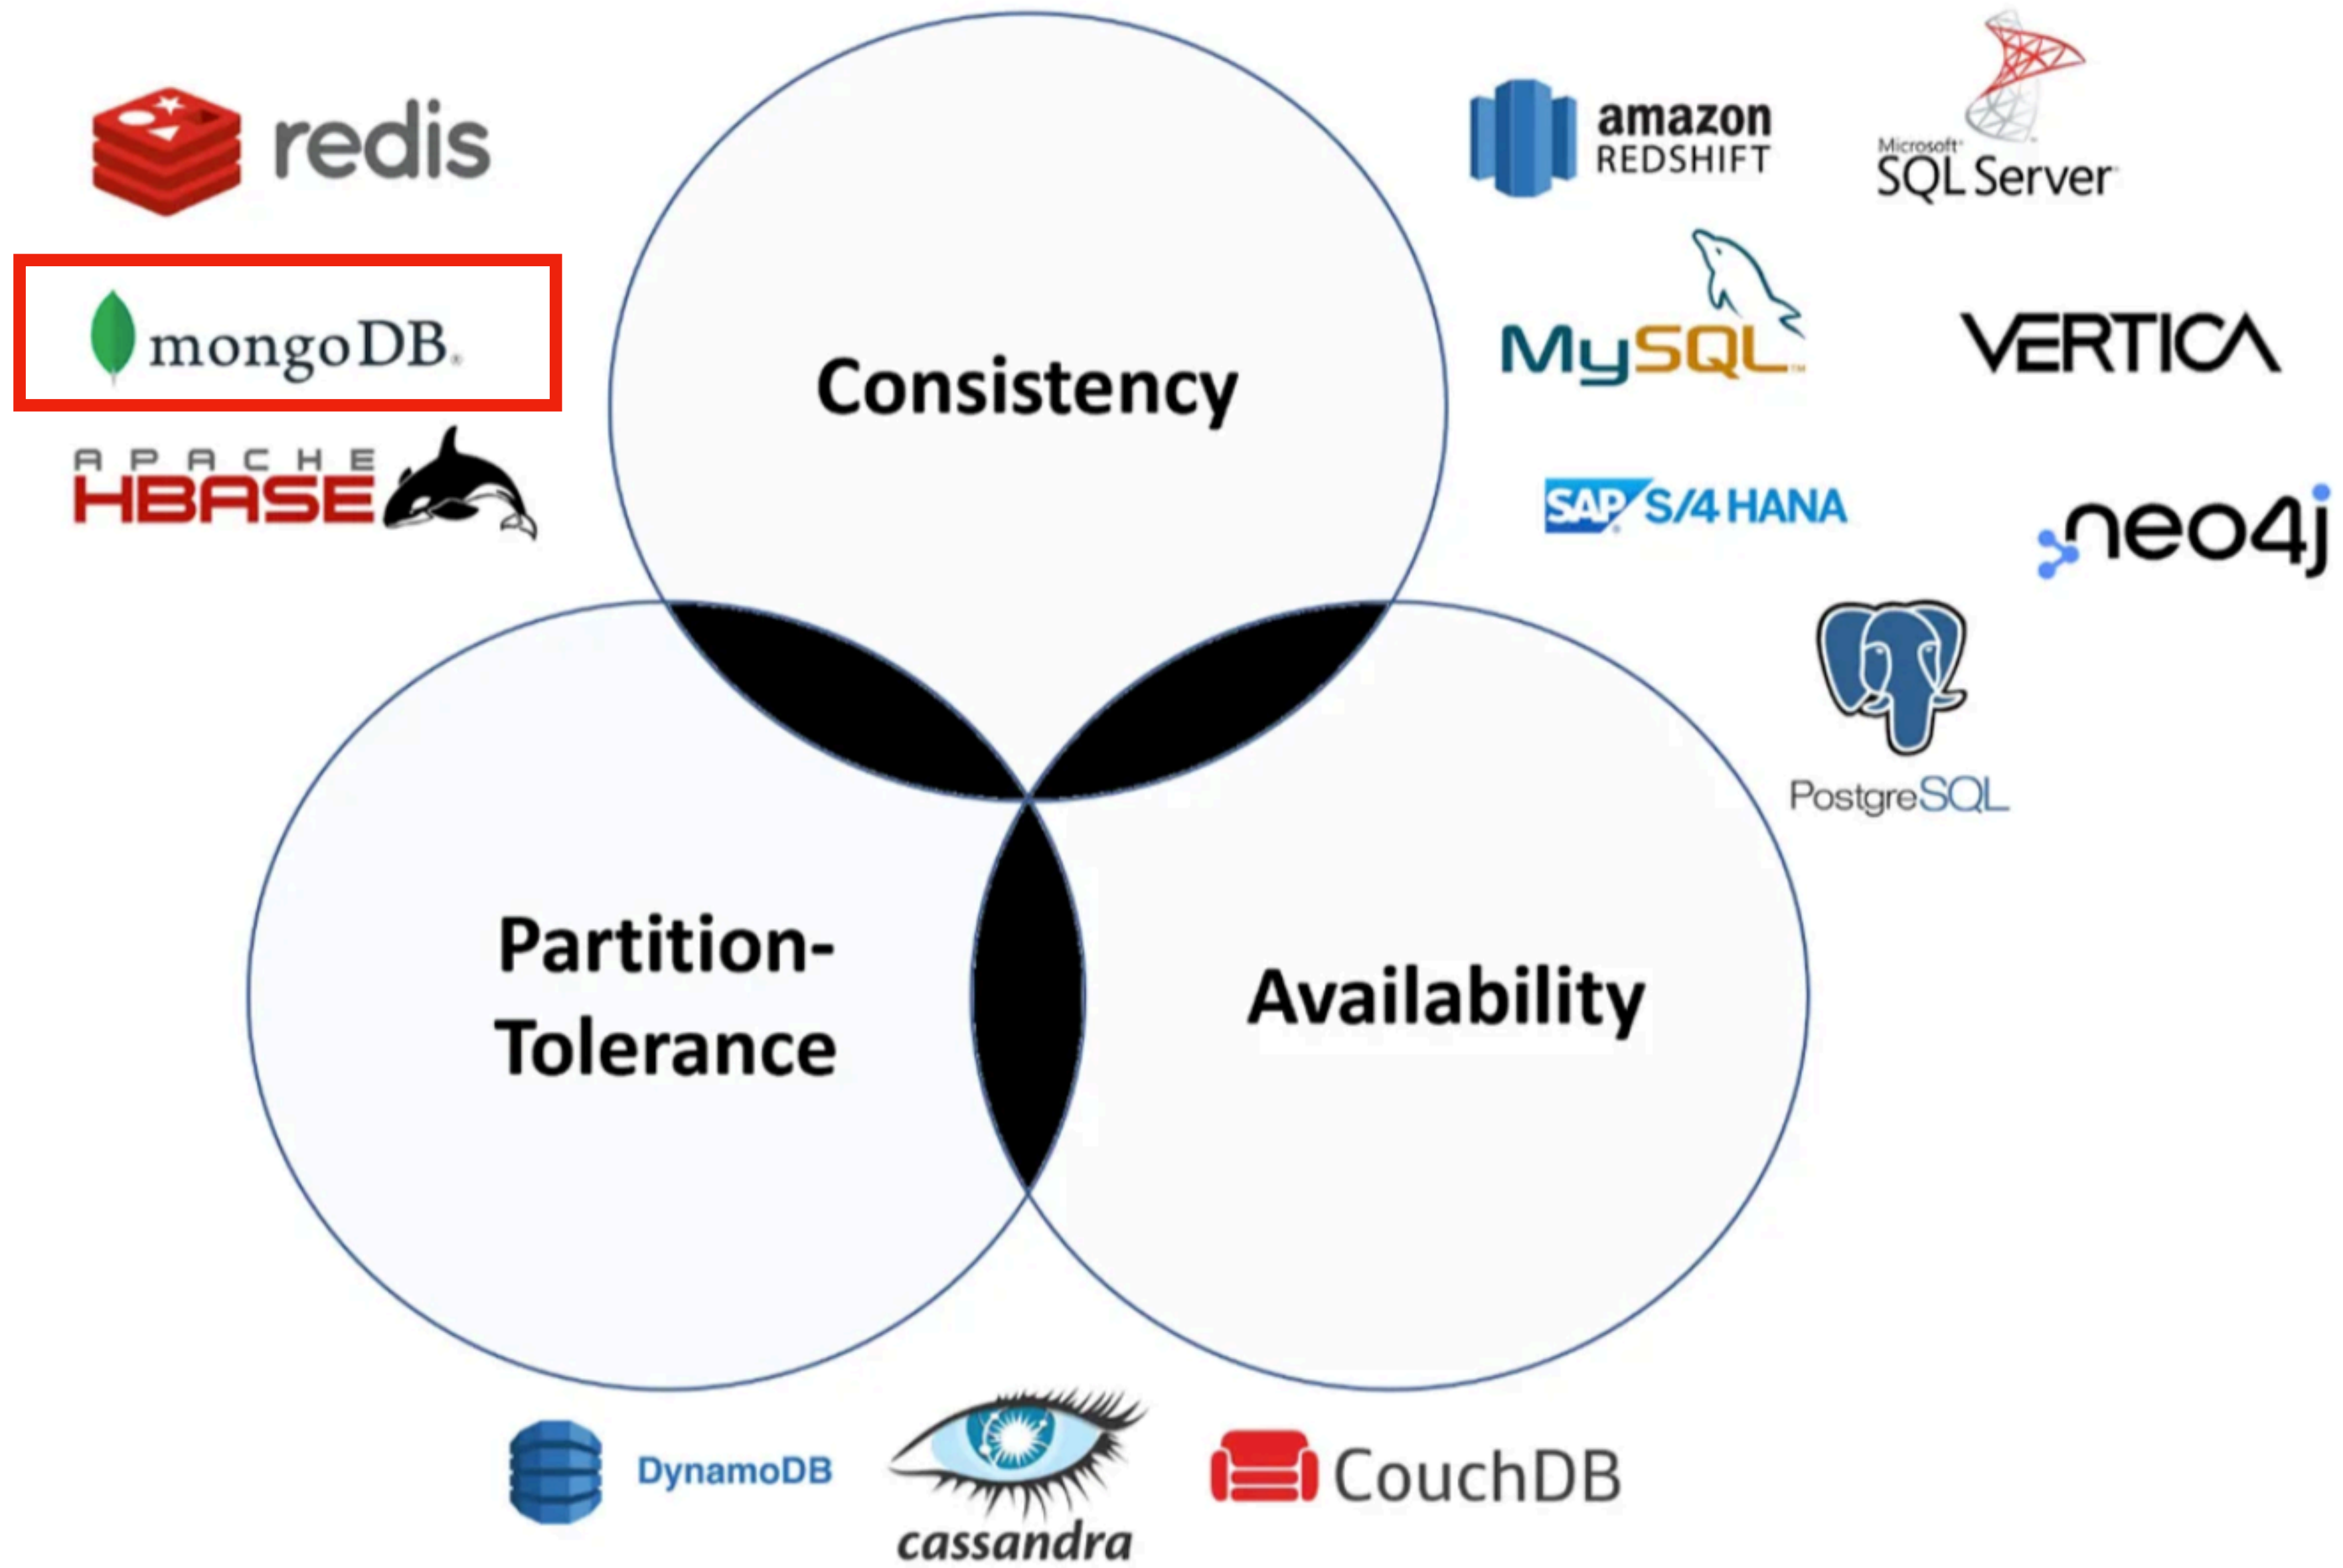
```
{
  "_id": 1,
  "first_name": "Leslie",
  "last_name": "Yepp",
  "cell": "8125552344",
  "city": "Pawnee",
  "hobbies": ["scrapbooking", "eating waffles", "working"]
}
```

- No need for joins
- No need for data normalization

When to use MongoDB?

Vs. MySQL

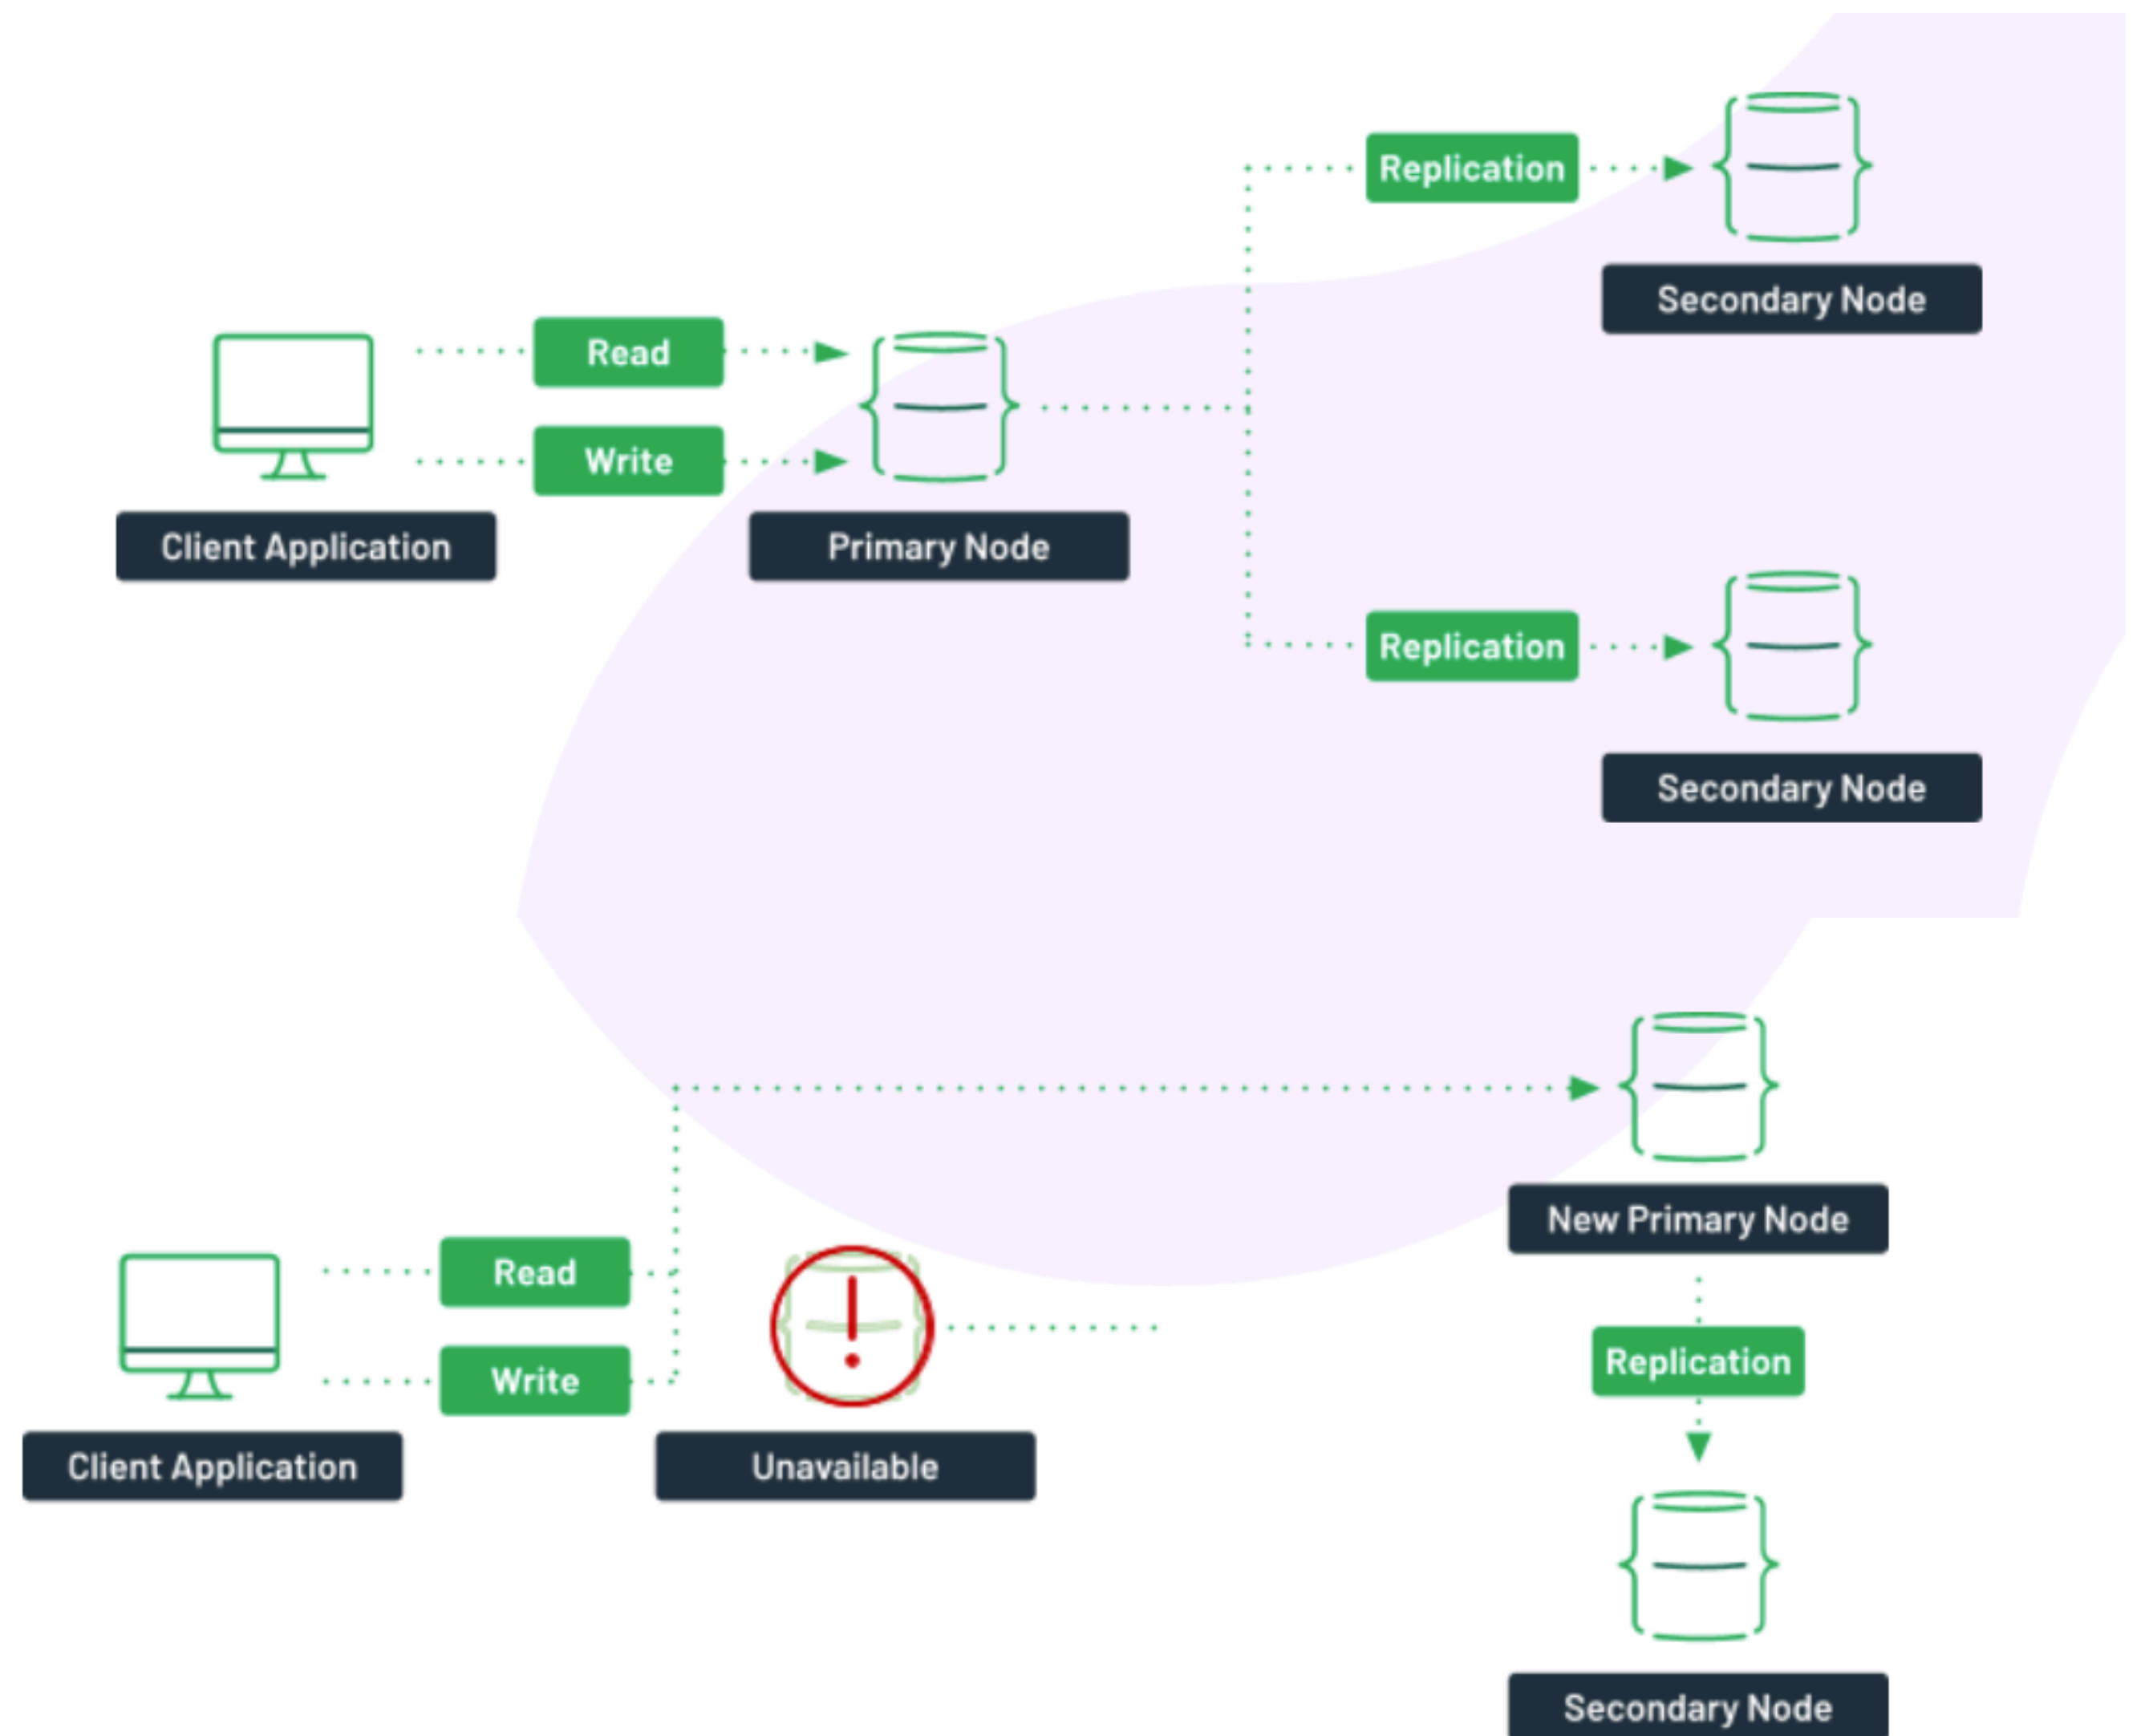
- Storage of unstructured and semi-structured data
- Huge volumes of data
- Requirements for scale-out architecture
- Modern application paradigms like microservice and real time streaming
- High write load
- Geospatial data
- Document-based data



MongoDB

Sacrifices availability

Split Brain Scenario




Installing MongoDB compass

Finding **fire** Pokemon with a
base_happiness > 50

Queries

Or Operator



```
{  
  $or: [  
    { "type1": "water" },  
    { "type1": "fire" }  
  ]  
}
```

Queries

And & OR Operator

```
$and: [  
  {  
    $or: [  
      { "type1": "grass" },  
      { "type1": "electric" }  
    ]  
  },  
  { "speed": { $gt: 80 } }  
]  
}
```

Queries

Less than

A code editor window with a light blue background and a grey title bar containing three colored window control buttons (red, yellow, green). The editor contains a JSON query snippet.

```
{  
  "base_total": { $lt: 300 }  
}
```

Queries

Regular Expressions

```
{
  $or: [
    // Names starting with "S" (case-insensitive)
    { "name": { $regex: "^S", $options: "i" } },
    // Names ending with "e" (case-insensitive)
    { "name": { $regex: "e$", $options: "i" } }
  ]
}
```

Spørgetime: Tid til at finde
spørgsmål